

Dragon

Decentralizing token data to discover alpha.

Co-created by Harshan and Frog

Last updated: June 1, 2025

email: baddragonmygoodness@gmail.com

Table of Contents

1 Introduction

2 Market Dynamics & Opportunity

3 Dragon

4 AlphaSwarm & Magic

5 Project Token

6 Development Goalposts

7 Founders' Note & Getting Involved

1

Introduction

Why aren't you in control of token analysis when you do your own research?

Dragon is a decentralized data platform powered by community-built token analyses. Each analysis produces specific knowledge on one aspect of a token project—sourced from either on-chain or off-chain data—and contributes to a growing library of token intelligence.

Today, the valuable token insights are scattered across platforms like Bubblemaps, BullX, Solscan, GMGN or locked behind Telegram bots, Discord servers, and paywalled research tools. Public insights from KOLs, Dune dashboards, and insider group chats are often unstructured and ephemeral. As a result, retail traders are forced to juggle dozens of tabs and subscriptions just to keep up. Discovering **alpha** becomes a losing game despite the fact that it's simply about collecting and interpreting quality data in real time.

Dragon gives any analyst the opportunity to contribute bite-sized, open-source **data modules** through a browser extension designed to streamline the DYOR (Do Your Own Research) process. Each user's Dragon is a customizable dashboard of token analyses, helping users discover alpha through their own intuitions and the collective intelligence of the community. With use, Dragon's AI trains on their strategic profile or "behavioral fingerprint," matching tokens they've bought to new tokens they haven't discovered yet.

As the platform grows, these analyses are ranked by impact and demand creating a composable, modular signal system for tracking shifts in volume, attention, and alpha. Dragon is more than a tool. It's an open infrastructure for building an intelligent layer of Web3, enabling anyone to impact the future of token research and development.

2

Market Dynamics & Opportunity

Trading bots have become a dominant force in crypto markets. Across blockchains, they've generated over \$80 billion in lifetime volume¹. On Solana alone, bot trading hit \$770 million in a single day in May 2025². As the fastest blockchain by transaction count³, Solana has become the launchpad for a new class of trading tools: token scanners.

Syrax, SarumAI, trench.bot, and rugcheck.xyz now serve hundreds of thousands of users per month⁴ providing automated safety heuristics, holder alerts, and social sentiment scores. These "DYOR assistants" are shaping how retail traders engage with memecoins and other high-risk, high-reward tokens. Even Solana-native trading platforms like Photon, BullX, and Axiom have integrated token scanning features directly alongside price charts, signaling that data transparency is now a core product expectation.

Despite this demand, token data remains fragmented and inconsistent. High-profile scandals⁵, extractive launches (e.g., pump.fun)⁶, and even obscure institutional behaviors⁷ have eroded trust, yet there's still no clear standard for retail-grade data transparency in token trading.

MEV & Information Asymmetry

One of the most misunderstood market dynamics in blockchain is Maximum Extractable Value (MEV), a design element that impacts roughly 25% of all transactions on Ethereum, Solana, and Binance Smart Chain⁸. During times of volatility, MEV can affect over 66% of block transactions⁹. It is frequently called the double-edged sword of crypto¹⁰: while it rewards validators and drives block production, it also creates information asymmetries that favor insiders and bots over everyday consumers. Infrastructure-level projects like Flashbots have emerged to address MEV on Ethereum, but there is no equivalent solution built for the retail layer; no way for traders to see how MEV might be affecting the tokens they interact with.

The Open Space for Retail Intelligence

So far, most Web3 data tooling has been focused on defense: compliance, threat detection, or top-down institutional reporting. Meanwhile, retail-facing analytics tools either lock insights behind paywalls or lack the flexibility needed to inform all traders' decisions.

Yet the market has proven their appetite. Dune Analytics validates the demand for community-driven dashboards. Nansen proves that trading data is immensely valuable especially when tailored to user cohorts.

What hasn't been built is a composable platform that combines the models: a community-driven, trader-focused, real-time token intelligence project that works for retail, by retail.

3

Dragon

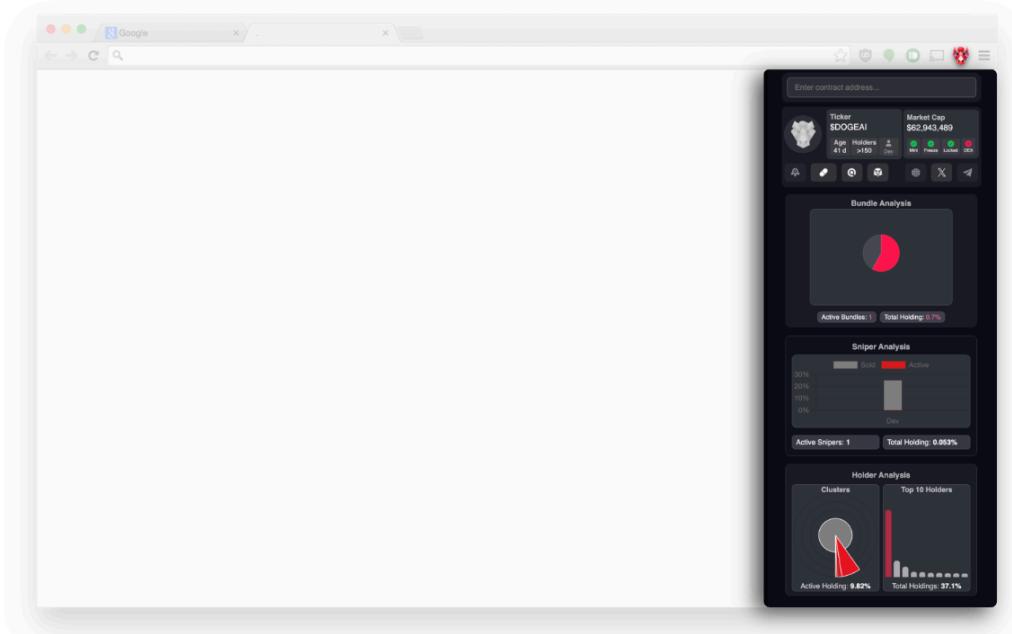


Figure 1: Dragon’s extension stays open inside your browser session, staying with you as you navigate the internet, trade tokens, or conduct DYOR.

To use Dragon’s browser extension, just paste in a Solana token’s contract address. The panel becomes a dashboard of real-time information on that token using community-developed analyses that draw from both on- and off-chain data. These analyses update and follow the user as they navigate the internet, becoming a consistent DYOR companion.

Dragon uses wallet-based single sign-on (SSO) to sync settings across devices and track user trading behavior. When a user purchases a token, Dragon snapshots the active analyses in their extension. Over time, this builds a behavioral profile that helps personalize recommendations and shape the trader’s strategic identity. By using read-only methods, Dragon keeps wallet security risks low while generating rich contextual data for the user.

Modularity

Dragon's key design feature is its modular architecture. Each token analysis is packaged as a bite-sized **data module** that evaluates one specific dimension of a token project—such as liquidity flow, smart money inflow, team wallet behavior, or developer history.

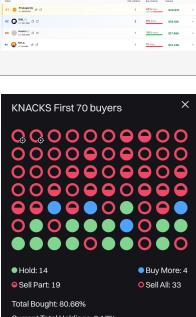
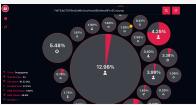
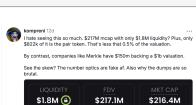
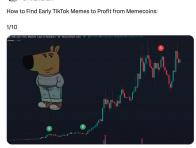
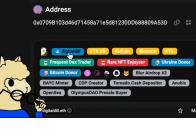
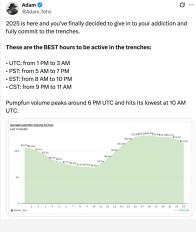
Dragon launches with four prototype modules, but the goal is to build an open ecosystem of community-contributed modules where anyone can propose and submit a new module that sourcing on- or off-chain data. If their data module is accepted, a developer earns a bounty in \$DRAGON to reward their contribution to retail trading intelligence.



*Figure 2: Token analyses are designed as customizable **data modules**—each focusing on a single token design metric. Users can mix, match, and even build these modules themselves.*

Inspiration

Token analyses today are fragmented among various individuals and different platforms. Dragon invites these data custodians into a shared, open-source research community where their insights can reach broader audiences, remain composable, and be continuously improved. By creating an arena for open competition, Dragon incentivizes better, faster, and more creative approaches to developing token research analyses.

Inspiration	Potential Analysis	Current Custodian	Inspiration	Potential Analysis	Current Custodian
	KOL / smart money holders	StalkChain		Token narrative profiling	Dexu / @cryptokoryo
	Sniper & insider buy counter	GMGN		Liquidity pool sandwich score	sandwiched.me
	JITO bundle visualizer	trench.bot		Top holder previous transaction behavior	RugChecker
	Top 10 holder awareness	Solscan		Dev / deployer overview	Trenchy Bot
	Clustered wallets	Bubblemaps		Holder trends by time held	HolderScan
	Liquidity - market cap ratio	@kompreni		Holder trends by time held	Birdeye (and others)
	CT / whale holder % and movement	@dethective		Holder profiles / style tags	DefiLlama
	TikTok / X sentiment matching	@PixOnChain		DEX paid? Project links and aesthetics	DEX Screener

Community Development

At launch, developers will propose new modules via pull requests on GitHub. These will be reviewed and integrated by Dragon's core team. As the project evolves, we will introduce a **Module Forge**—a web-based drag-and-drop sandbox where developers (including non-JavaScript-native users) can build modules in Python or Rust, compile them to WebAssembly (WASM), and test them against historical token datasets in real time.

A complementary **Module Market** will allow users to upvote or downvote modules, request new modules to be developed, and toggle modules on/off instantly without needing to reload their extension. This creates a fully composable and permissionless platform where community demand drives module development, visibility, and utility.

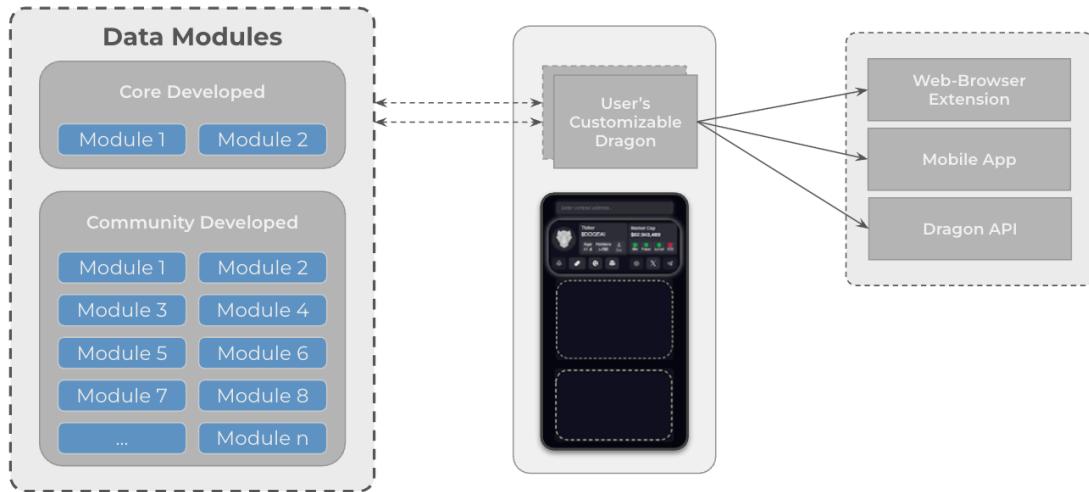


Figure 3: Dragon's architecture includes both core developed modules and community developed modules. The end product syncs across a user's browser extension, a mobile app, and a public API.

In time modules will be ranked by their performance, based on their correlation to real token volume changes and their adoption across Dragon users. High-impact community developers will gain more influence in the project's direction, aligning incentives between module quality and community value.

Mobile Application

Dragon's mobile app mirrors the extension, giving users access to their modules, dashboards, and alerts on the go. Push notifications can be configured to trigger when selected modules detect major shifts such as price anomalies, social spikes, or liquidity surges.

Because Dragon tracks usage patterns across both desktop and mobile, its personalization engine (powered by the Magic engine) can adjust for behavior deltas between devices, refining token match opportunities to traders where they are.

4

AlphaSwarm & Magic

An orchestration agentic approach to selective LLM training that captures Web3 alpha with a user's modules, sourcing on- and off-chain data and historical volume changes.

Introduction

AlphaSwarm is a multi-agent observatory designed to recognize token movements by fusing module outcomes, turning on-chain events (DEX volume, wallet flows, MEV activity) and off-chain signals (social sentiment, news, alpha groups) into actionable trading insights. The system doesn't rely on full-streaming data. Instead, it learns from significant market moments—volume spikes, rug pulls, pumps—and builds a behavioral map of modular signal combinations during those key events.

These snapshots of module data and volume fluctuations become training data for a central LLM-powered orchestrator agent that learns trend patterns holistically, rather than isolating the effect of a single signal. It rates new events on a 0–10 scale (quantitative) based on how much they resemble past scenarios and provides a narrative explanation (qualitative)—offloading the cognitive burden from traders.

Event-Curated LLM Training on Meaningful Windows

Rather than training the Large Language Model (LLM) on an exhaustive stream of real-time data—much of which is noisy or uninformative—Dragon adopts a more efficient and insightful approach: event-curated training. Specifically, the system identifies periods in the dataset with significant volume changes and extracts the full set of module outputs corresponding to those timeframes. These time-bounded, high-impact segments—referred to as event windows—serve as the training corpus for the orchestration LLM.

Each data module in Dragon is handled by a dedicated agent that monitors and logs its respective features over time. The LLM then learns how different combinations and permutations of these module trends correlate with actual market movements. Rather than attempting to assign linear causality to any single feature, the LLM learns complex, multi-dimensional patterns that characterize specific market scenarios. During inference, when real-time data is fed into the system, the orchestrator evaluates the new inputs against the

historical pattern library and computes a similarity score to past events. This methodology enables scenario recognition based on learned multi-signal patterns, offering a more robust and context-aware form of predictive analytics.

Agentic Confounding Resolution via Orchestration

One of the fundamental challenges in financial modeling is the issue of confounding, where multiple variables change simultaneously, making it difficult to isolate the effect of any single factor. AlphaSwarm avoids the faulty attempt of strict real-time causality per module, which is typically unreliable in the stochastic and noisy environment of crypto markets. Instead, the system relies on a team of agentic modules, each tasked with monitoring a specific metric—such as **Sniper %** or **Bundle %** in a token—individually over time.

These agents feed their observations into the orchestration LLM, which is trained to evaluate the collective context of module behaviors rather than individual contributors. For example, if past event data shows that sharp increases in **Sniper %** combined with flat **Bundle %** and rising liquidity often lead to significant frontrun-pump scenarios, the model learns to recognize that pattern. When similar conditions reappear in live data, the orchestrator surfaces the match—not as a hard statistical claim, but as a pattern-aligned scenario based on prior outcomes. In this way, the system resolves confounding, using coordination and historical pattern matching instead of assuming modular independence.

Qualitative & Quantitative Output Layer

AlphaSwarm delivers a scenario-based scoring mechanism ranging from 0 to 10. A score of 0 indicates that current conditions bear no resemblance to known high-volatility scenarios from the training data, whereas a 10 reflects a strong historical match to a past pattern with significant market impact. This scoring system reflects confidence in the model's learned memory rather than a strict statistical likelihood.

Each score is accompanied by a breakdown of per module historical precedence—indicating which components contributed most to the pattern match—and a natural language explanation. The explanation is generated through an embedded NLG (Natural Language Generation) pipeline, allowing traders to quickly interpret the rationale behind each signal. This dual output of numeric score and qualitative insight helps traders build intuition and make informed changes in their trading focus, even in situations where pure mathematical certainty is lacking.

AlphaSwarm Architecture

The AlphaSwarm is built as a robust, real-time data analytics system that integrates a full-stack modular AI pipeline. It sources data from the various community developed modules that source both on-chain and off-chain events.

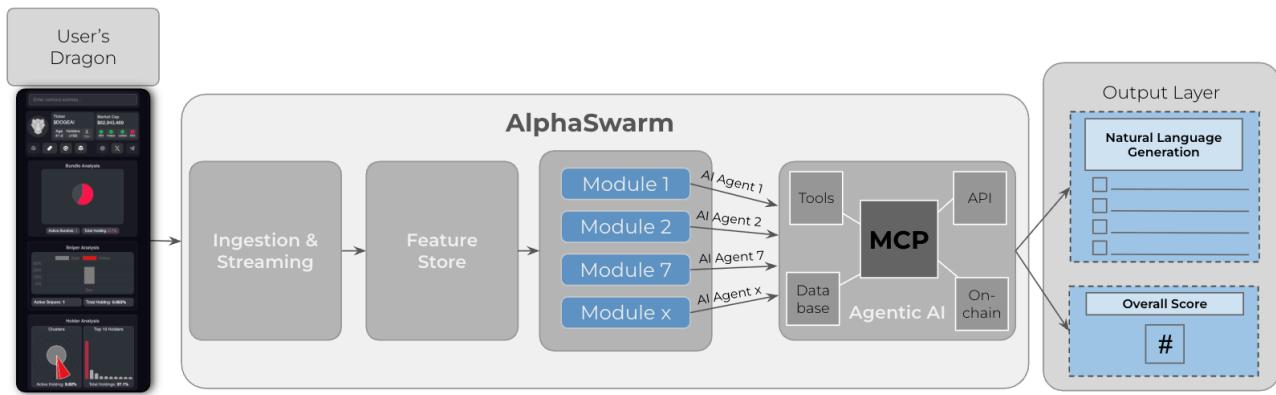


Figure 4: Dragon's Core AlphaSwarm Module

All of the data streams are ingested through a high-throughput infrastructure using technologies like AWS Kinesis or Apache Kafka, which organize the signals into time-aligned buffers. These are then written into a dual-layer feature store: a hot layer (e.g., Redis or KeyDB) for sub-second latency access and a cold layer (e.g., Apache Iceberg on S3) for long-term historical analytics and backtesting.

The core of AlphaSwarm is the modular analysis stack, composed of discrete agents working on each module. These module outputs are passed to the Agentic AI hub, or **Model Context Protocol (MCP)**, where agents using LangGraph or AutoGen frameworks validate, enrich, or challenge one another's findings to orchestrate the final output.

What is Model Context Protocol (MCP)?

MCP is an open standard, open-source framework introduced by Anthropic to standardize the way LLMs integrate and share data with external tools, systems, and data sources. Inspired by agentic AI systems like Microsoft's AutoGen, MCP ensures real-time adaptability without retraining the full pipeline, echoing a broader shift in corporate AI towards modular, context-aware intelligence.

For more information on MCP, visit Anthropic's announcement [here](#).

In Dragon, the MCP is a coordination layer that dynamically assigns trust scores to analytic modules based on how well their signals align with actual volume changes. It acts like an ensemble controller—boosting modules with strong predictive value. The MCP serves as a shared trust layer, dynamically updating the weight w_i of each module based on its historical signal quality. A final composite hype index $H = \sum_i(w_i \times signal_i)$ is computed for each token, where:

$$H = \text{total hype score}$$

$$\sum_i = \text{sum over modules}$$

$$w_i = \text{module trust weight}$$

$$signal_i = \text{module output signal}$$

AlphaSwarm's results are rendered in two formats through its output layer: a quantitative score from 0-10 (and perhaps a future radar chart!) that summarizes the weighted module contributions and a natural language summary that explains the reasons why a token may have a significant shift in volume soon. These outputs are delivered via the AlphaSwarm API to both the browser extension and mobile app, offering an immediate and accessible interface for traders. The entire system is also integrated with the **Magic** engine, enabling one-click new token discovery based on a user's trading patterns.

A Personalized Magic Module

Dragon enhances AlphaSwarm's framework with **Magic**— a personalized module designed to increase your action potential as a trader. Every time you interact with another data module or make an on-chain transaction, your behavior is captured and embedded into a dynamic user profile. This profile encodes not just your preferences, but your actual trading behaviors—the risk levels, liquidity patterns, and market sentiment signals at which you engage with tokens at the most—forming a “behavioral fingerprint” of your strategic profile based on the data analyses that you use.

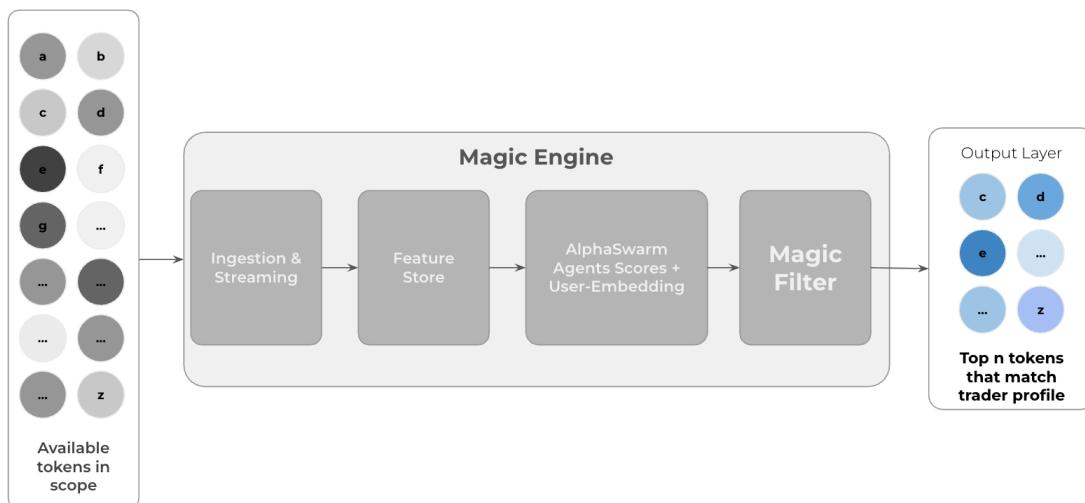


Figure 5: Dragon's Core Magic Module

This fingerprint becomes the input to a two-tower transformer architecture at the core of the Magic engine. The first tower ingests your usage history: module activation patterns, past transactions, and clickstream events across the extension. The second tower encodes AlphaSwarm's fused analytics that have been passed through the trust-weighted MCP. This means every token embedding already reflects the most relevant AlphaSwarm-derived signals.

During training, the system optimizes a learning objective that prioritizes tokens you've historically interacted with or traded. Over time, your preferences become more refined and context-aware as AlphaSwarm agents re-weight signal quality, ensuring that your new tokens are not just personalized, but also responsive to changing market dynamics.

For fast and scalable retrieval, Dragon indexes all token embeddings using FAISS's HNSW (Hierarchical Navigable Small World) algorithm, enabling near-instantaneous nearest-neighbor lookups. Your personal embedding is updated in real-time using lightweight online learning steps and consolidated with a nightly batch retrain to capture longer-term evolution. For new users without enough interaction history, the system applies clustering algorithms like DBSCAN to assign a cold-start profile by grouping you with traders showing similar early-stage signals.

The Magic Button Experience

Once the Magic engine identifies a set of top-matching tokens based on a fusion of your behavioral fingerprint and AlphaSwarm's signals, the Magic button offers a one-click bridge to execution. With a single tap, the Magic Button filters out all non-matched tokens on your selected platform (e.g., DEX Screener, pump.fun), and preloads parameters. In effect, AlphaSwarm fuels the intelligence, Magic delivers the personalization, and the Magic Button enables an action.

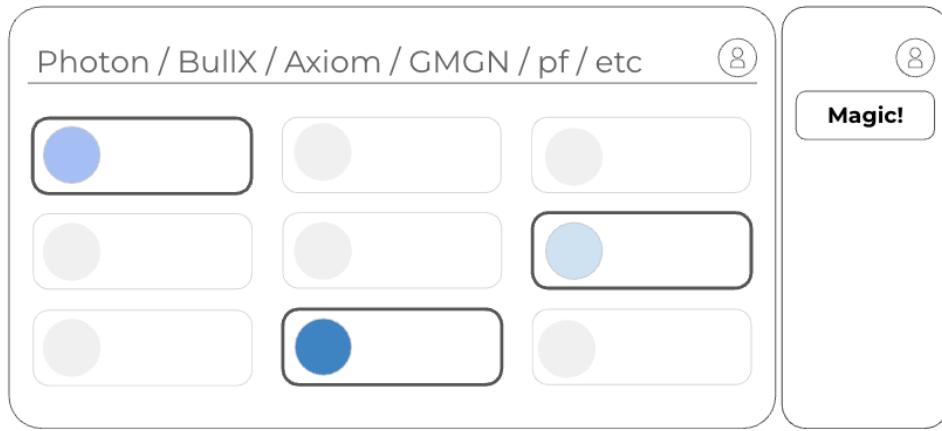


Figure 6: The Magic button interacts with your browser, eliminating the tokens onscreen that do not match your trading profile.

This integration of insights, preferences, and execution expedites the loop for intelligent trading—giving users an advantage during new token cycle iterations.

Parallel Enhancements to AlphaSwarm & Magic

Dragon will eventually include a built-in conversational chatbot—your on-demand trading assistant embedded in the extension and mobile app. This chatbot will leverage a retrieval-augmented generation (RAG) pipeline to pull in exactly the right Wikipedia passages, module data, and real-time agentic analytics you ask for. LangChain agents will orchestrate tool calls for signal queries, while Dragon connects to your encrypted profile and transaction history so every answer is personally tailored—becoming your true AI companion.

5

Project Token

CA:

Utility

There are three utilities the \$DRAGON token provides to holders.

1. Rewards a developer for a successful contribution of a token analysis module.
2. Enables a user's access to the core AI modules (AlphaSwarm and Magic).
3. Facilitates a DAO governance of the project's development and treasury allocation.

These three functions will be enabled gradually and with increased detail as the project progresses.

Rollout

- I. \$DRAGON is launched using Believe: 1% of all trading fees go to the project's treasury, 0.5% in SOL and 0.5% in \$DRAGON
- II. Project contributors are given tokens from the \$DRAGON treasury to promote community growth and project participation
- III. Project overhead and administrative costs are paid out of the SOL treasury to streamline transactions and allow for the widest scope of development opportunities
- IV. When trading fees and project growth have stabilized, the core team will build a smart-contract that automates treasury distributions and payments
- V. When ready, the smart-contract is handed over to a DAO of \$DRAGON token holders to decentralize control of the treasury and the project's development
- VI. The DAO and core team keeps building together, fully decentralized

The governance rules and logistics of the DAO will be decided in time, with both the token holders and the core team's combined input.

6

Development Goalposts

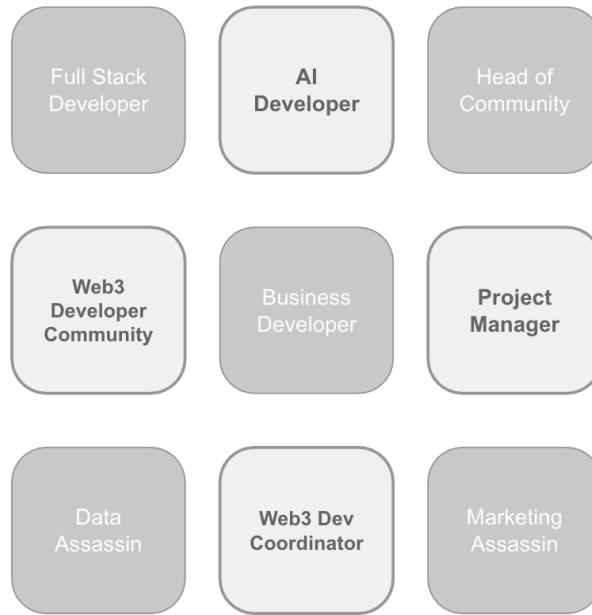
June	Team hired!	Wallet login enabled	Treasury automated and justified
July	Charts redesign for first modules	RPCs built to first modules	Logo & mascot redesigned
Aug	Customizable extension	Dataset created: AlphaSwarm v0	Web hub for customization
Sep	Training begins: AlphaSwarm v0	Module Forge developed	
Oct	Training begins: Magic v0	First release: AlphaSwarm v0	DAO infra begins
Nov	Mobile app	Magic v0 first release	Multi-database integrations
Dec		ETH / BSC / ... expansions?	DAO takes the wheel

[^]Dependent on labor & resources. Could go faster if we go higher and slower if we go lower.

7

Founders' Note & Getting Involved

This project has been launched by Harshan ([Harshit](#)) and Frog ([Andrew](#)), an AI researcher and a project manager by background. We consider Dragon's third founder to be the Web3 developer community, which wasn't the original plan but neither of us brings blockchain development experience to the table. Fortunately this became a strong feature of the business design as a community-powered developer platform, where open modular contributions create market-valued innovations for product users and traders.



*If you are interested in joining the team, please send an email to
baddragonmygoodness@gmail.com with an introduction and your relevant experience!*

Our goal now is to build a lean and committed core team, including a Web3 developer to coordinate community module integrations. A core tenet of the project is how connected it is to market demand. If it turns out that modular and open-sourced token data is popular among retail traders, we are prepared for the opportunity to publish research on user data as it relates to transactions, volume, and effective token analysis modules. The insights gained from this research promises value for blockchains, exchanges, and firms invested in better token designs and the intelligent expansion of Web3.

Our advantage lies in our daily exposure to institutional-grade AI and research, which we apply directly to this project's development. Harshit and Andrew met in a Masters of Data Science program at the University of Maryland at a time when Harshit had just completed research at IBM [co-developing in-house generative AI systems](#). Harshit currently works at Amazon building scalable infrastructure for Bedrock AI and developing the latest implementations of MCP methodology. Andrew approached Harshit for tutoring in the program, but ended up trading memecoins during sessions instead of learning the material. He previously co-founded the blockchain startup Cent, which raised a seed round after their [Valuables project gained attention for tokenizing tweets](#) as NFTs on Ethereum. Andrew brings a multi-disciplinary perspective to project management, drawing from an international background and experiences managing global campaigns and creative projects outside of the tech bubble.

We see ourselves as outsiders jumping in the deep end—retail memecoin traders that lost money to sniper bots and crypto-insiders. We look forward to both stewarding and being stewarded by a vision-forward community that shares our mission: to equip everyday traders with data-driven alpha through open intelligence.

Contribute to Dragon

If you're a developer, analyst, or Web3 native who likes the idea of open-sourcing alpha, we invite you to co-design this project's foundation with us. We're on the ground floor.

Build a module

Submit your idea to [GitHub](#) for an on- or off-chain analysis that surfaces real-time token insights. This can be written in Python, Rust, or JavaScript and eventually compiles to WebAssembly via our Module Forge (coming soon). Bounties in \$DRAGON are awarded for modules accepted to the platform.

Join the community

Connect with us and other contributors in our [Telegram](#) to share ideas, plan development, and discuss upcoming features and module priorities.

Collaborate on research

We're looking to collaborate with researchers on topics like token behavior modeling, retail trading patterns, and design ideas for high-impact analyses. Reach out to [Andrew](#) directly on Telegram for partnership opportunity.

Join the core team

We're currently looking for a few experienced assassins, including a Web3 dev to lead module coordination and infrastructure development. Send us an [email](#) or send [Andrew](#) a message on Telegram if that's you. More open positions will be posted on our [X](#) account!

Project Links

Website

<https://alpha-dragon.ai/>

Extension Prototype

<https://chromewebstore.google.com/detail/dragon/ncbglgbplhnbekllhogabdefjedbkoee>

GitHub

<https://github.com/alpha-dragon-org>

X / Twitter

<https://x.com/AlphaDragonAI>

Telegram

<https://t.me/+OU0SLVfcEZhZWQx>

Demo

<https://vimeo.com/1062123553>